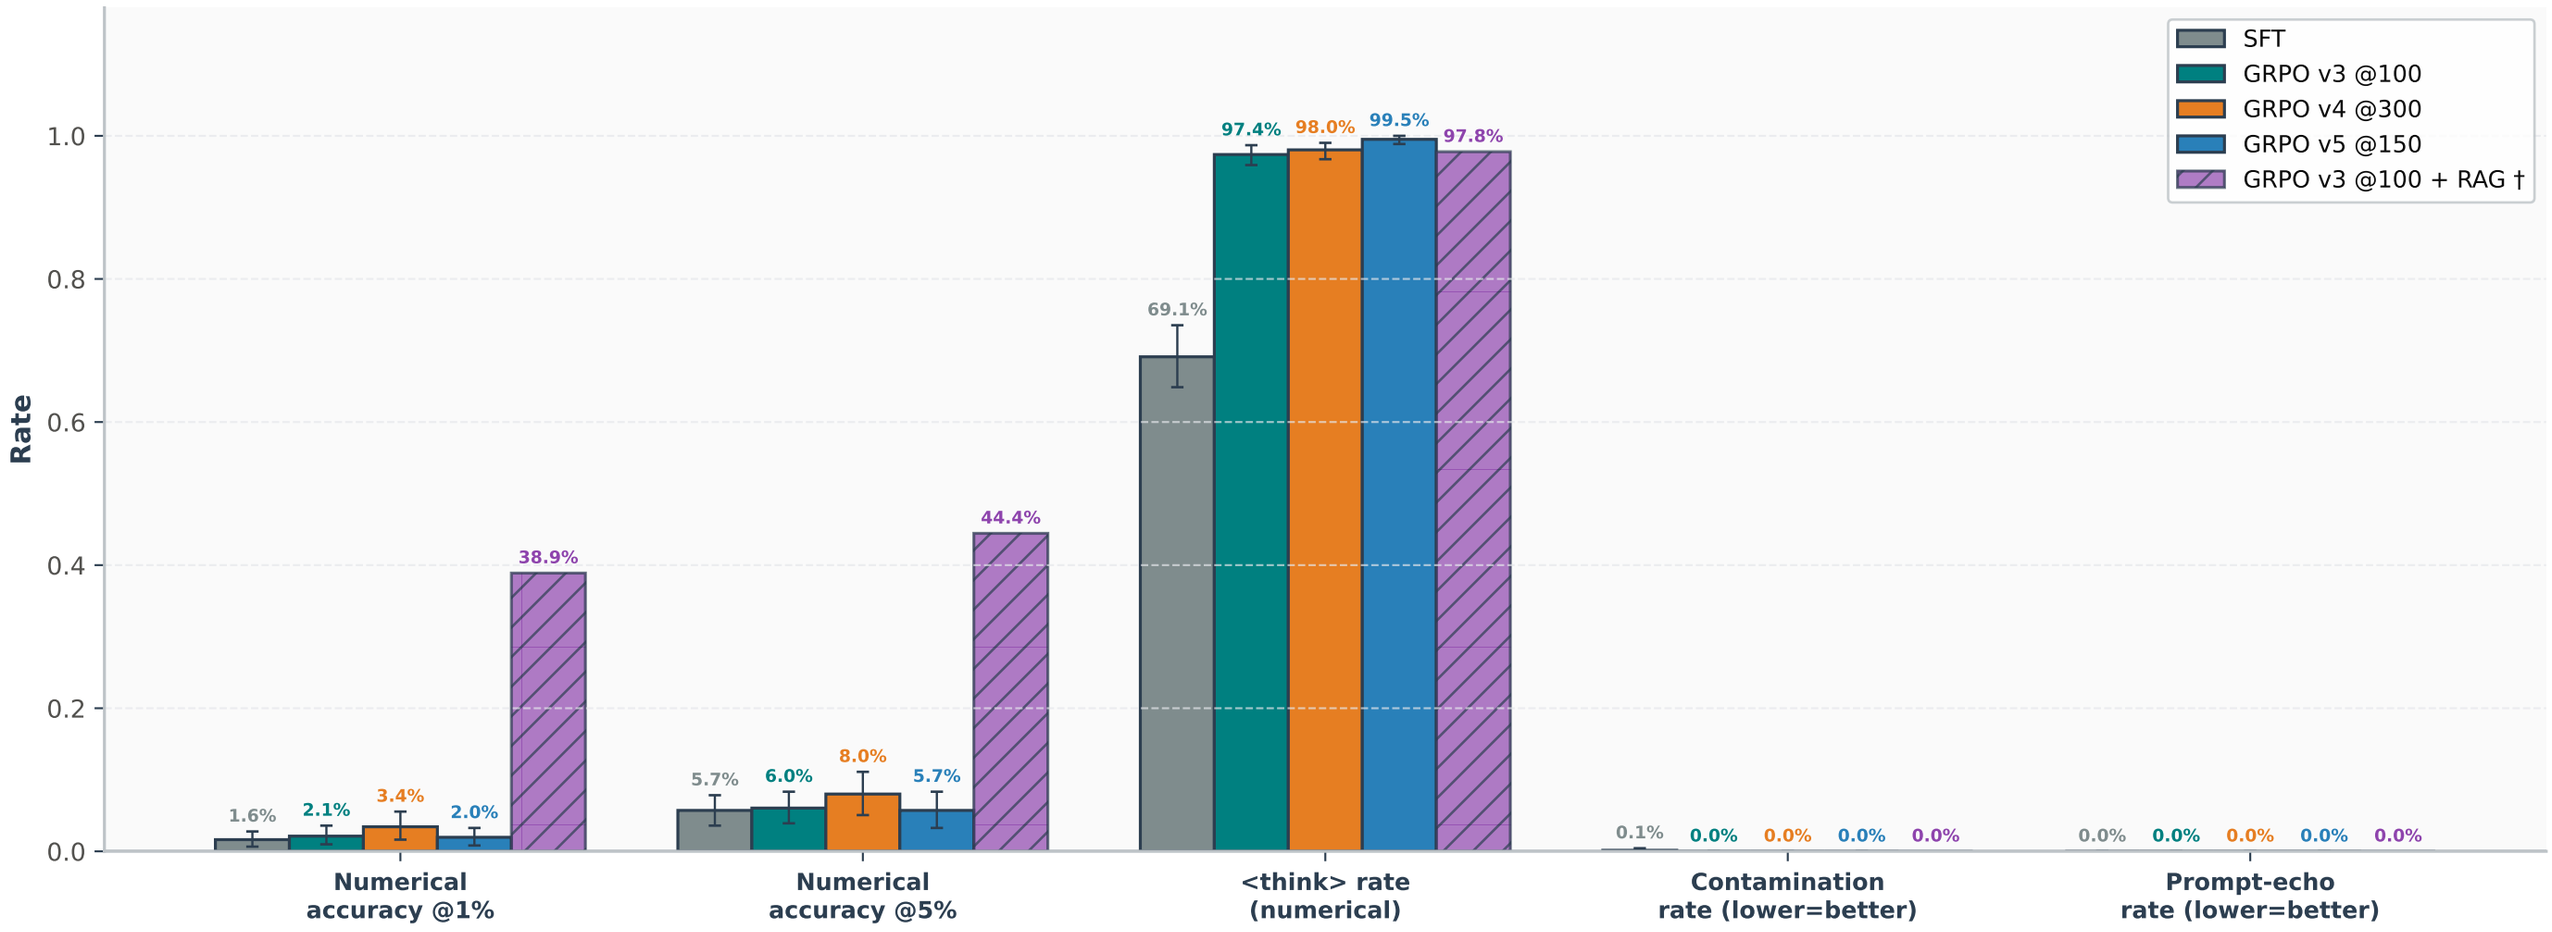


*Solid bars: 153 numerical + 20 conceptual held-out prompts × 4 samples, corrected (v2) labels, PITFALLS C3; pinned decoder (temp 0.3 / top-k 40 / rp 1.2); whiskers = 95% bootstrap CI over prompts. [eval_rl_v2labels_k4.json]*
*v5 vs v4 here (5.7% vs 8.0%) is one seed set; over 8 seeds x 4 samples with the batched decoder they are equal (6.4% vs 6.4%, docs/eval_vote_pooled8.json).*
*† Archived open-book run (docs/eval_rl_results.json): a different population (90 test-half numerical prompts) scored on the old (v1) labels; no committed script reproduces it.*