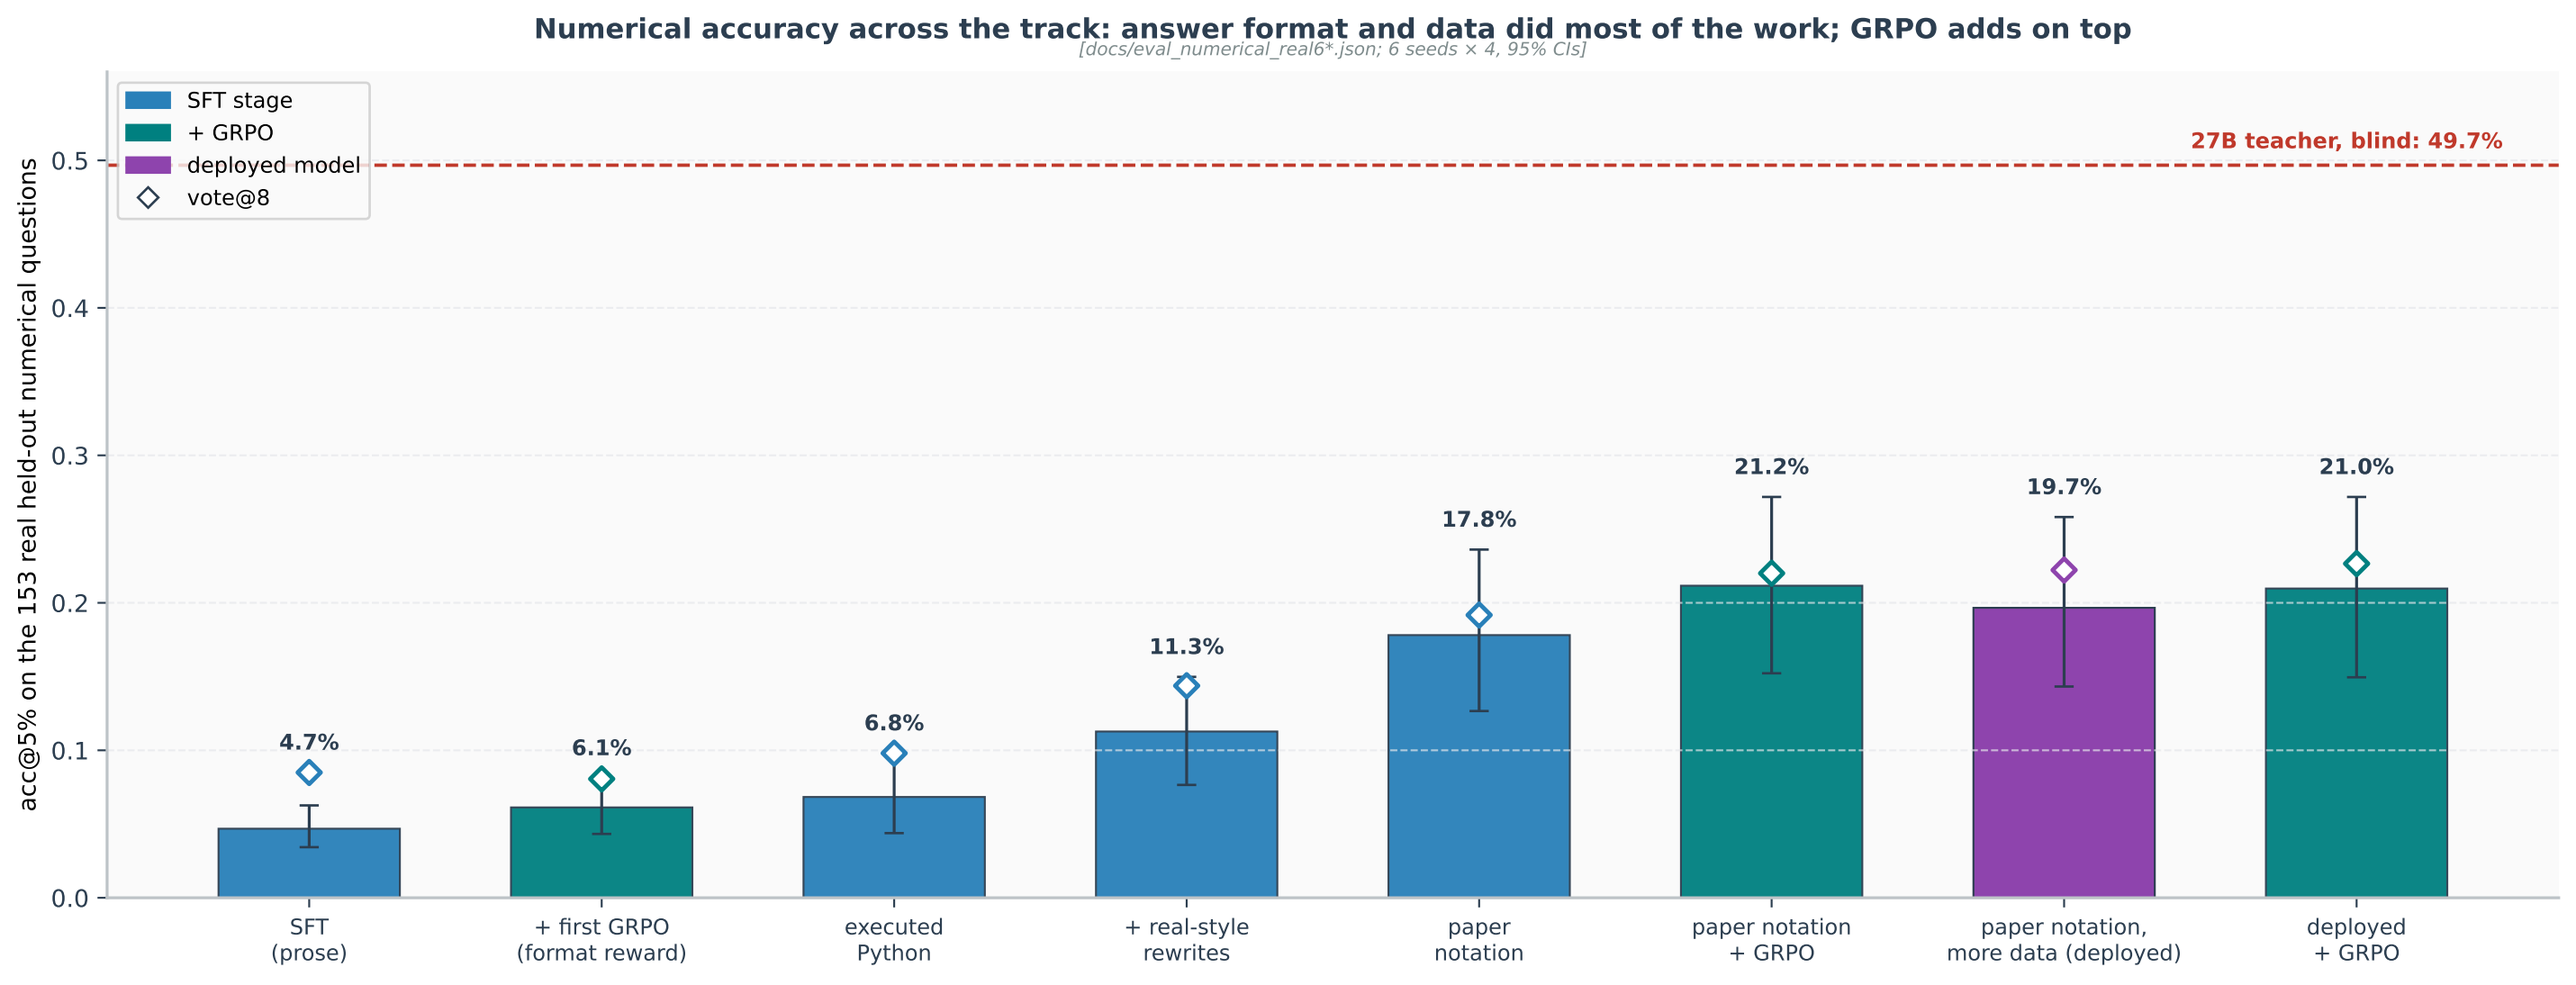
Numerical accuracy across the track: answer format and data did most of the work; GRPO adds on top
[docs/eval_numerical_real6*.json; 6 seeds × 4, 95% CIs]
SFT stage
+ GRPO
deployed model
vote@8
27B teacher, blind: 49.7%
acc@5% on the 153 real held-out numerical questions
0.0
0.1
0.2
0.3
0.4
0.5
4.7%
6.1%
6.8%
11.3%
17.8%
21.2%
19.7%
21.0%
SFT (prose)
+ first GRPO (format reward)
executed Python
+ real-style rewrites
paper notation
paper notation + GRPO
paper notation, more data (deployed)
deployed + GRPO