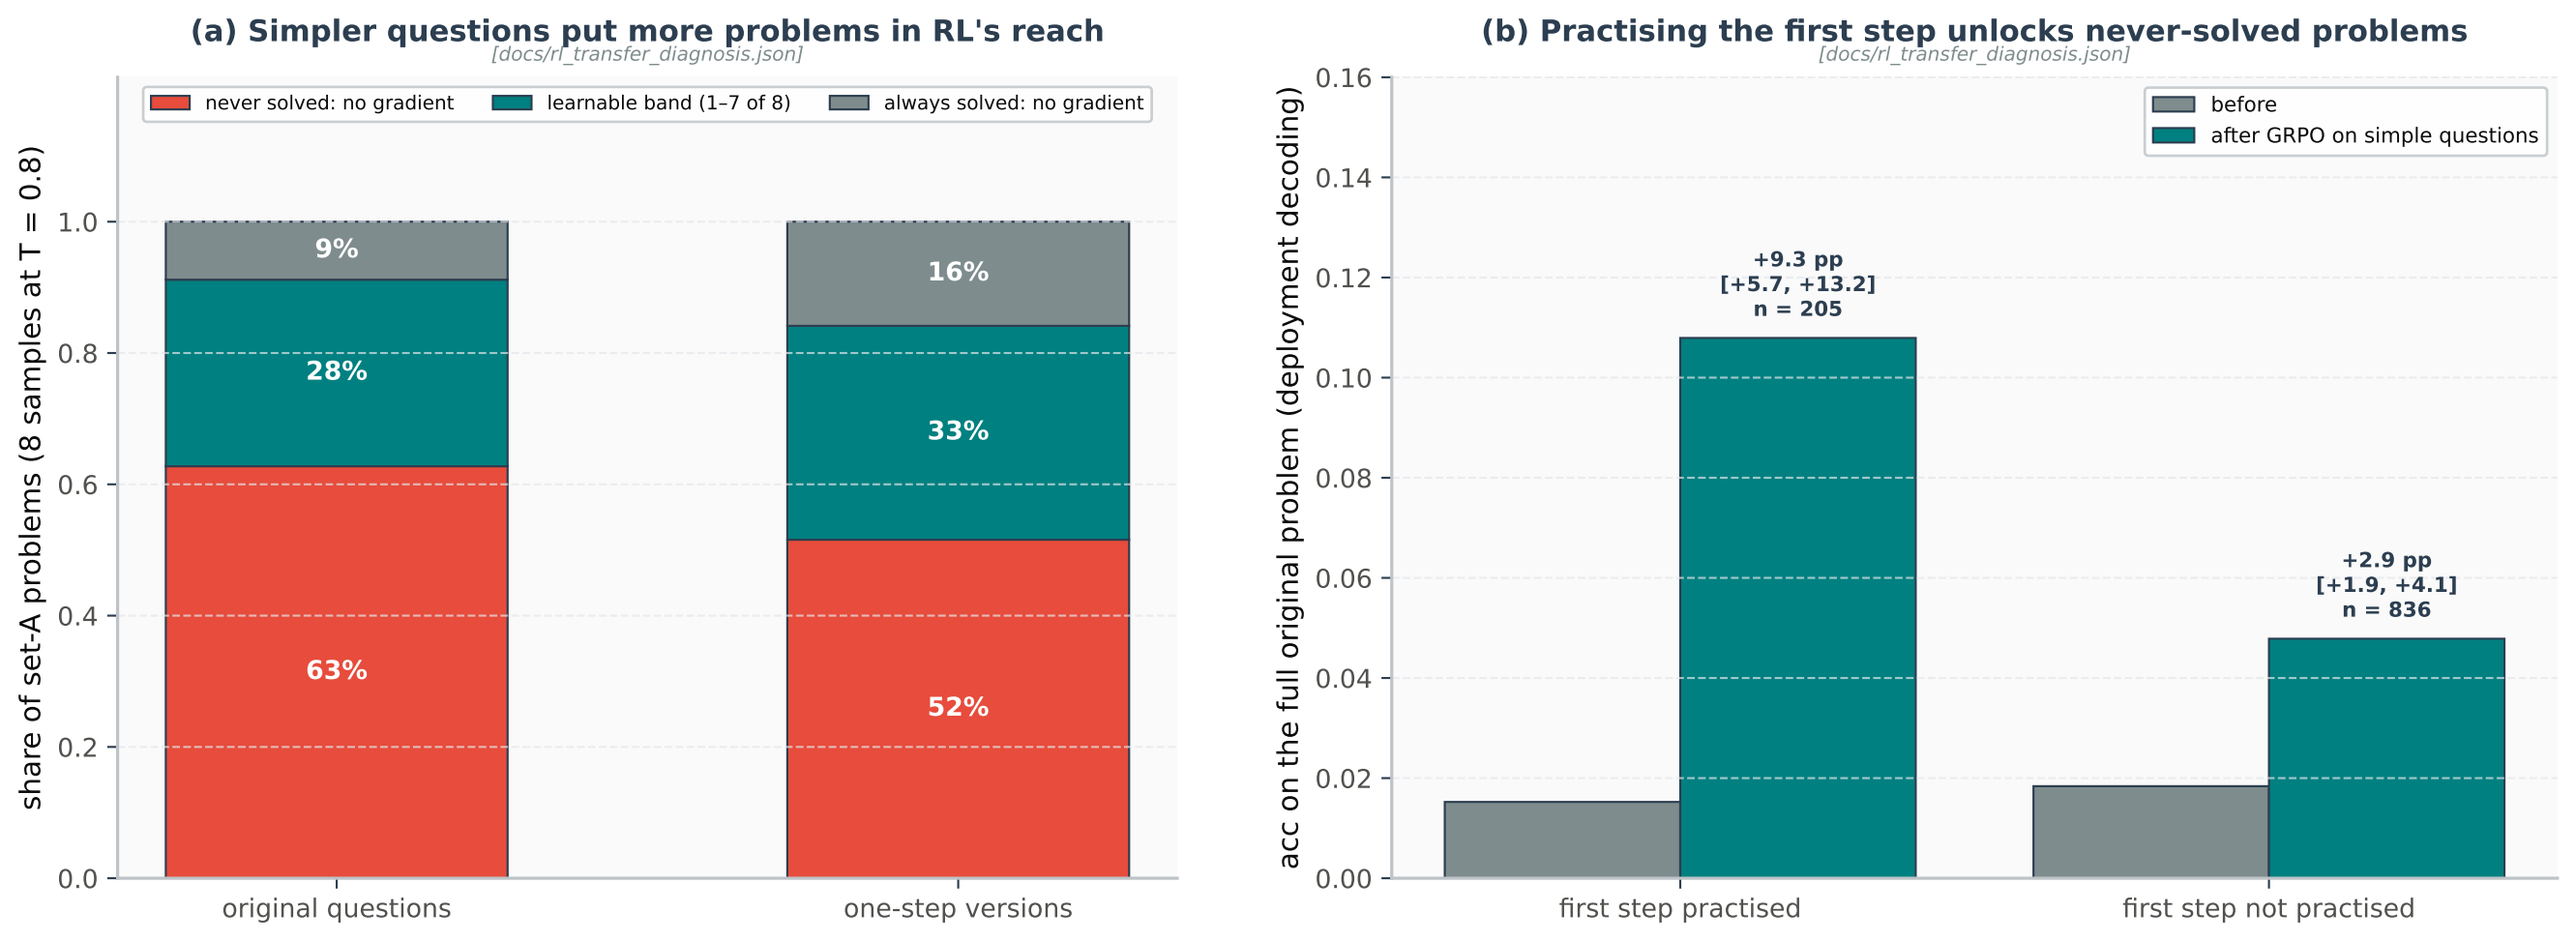
(a) Simpler questions put more problems in RL's reach
[docs/rl_transfer_diagnosis.json]
never solved: no gradient
learnable band (1–7 of 8)
always solved: no gradient
share of set-A problems (8 samples at T = 0.8)
0.0
0.2
0.4
0.6
0.8
1.0
9%
28%
63%
16%
33%
52%
original questions
one-step versions
(b) Practising the first step unlocks never-solved problems
[docs/rl_transfer_diagnosis.json]
before
after GRPO on simple questions
acc on the full original problem (deployment decoding)
0.00
0.02
0.04
0.06
0.08
0.10
0.12
0.14
0.16
+9.3 pp
[+5.7, +13.2]
n = 205
+2.9 pp
[+1.9, +4.1]
n = 836
first step practised
first step not practised