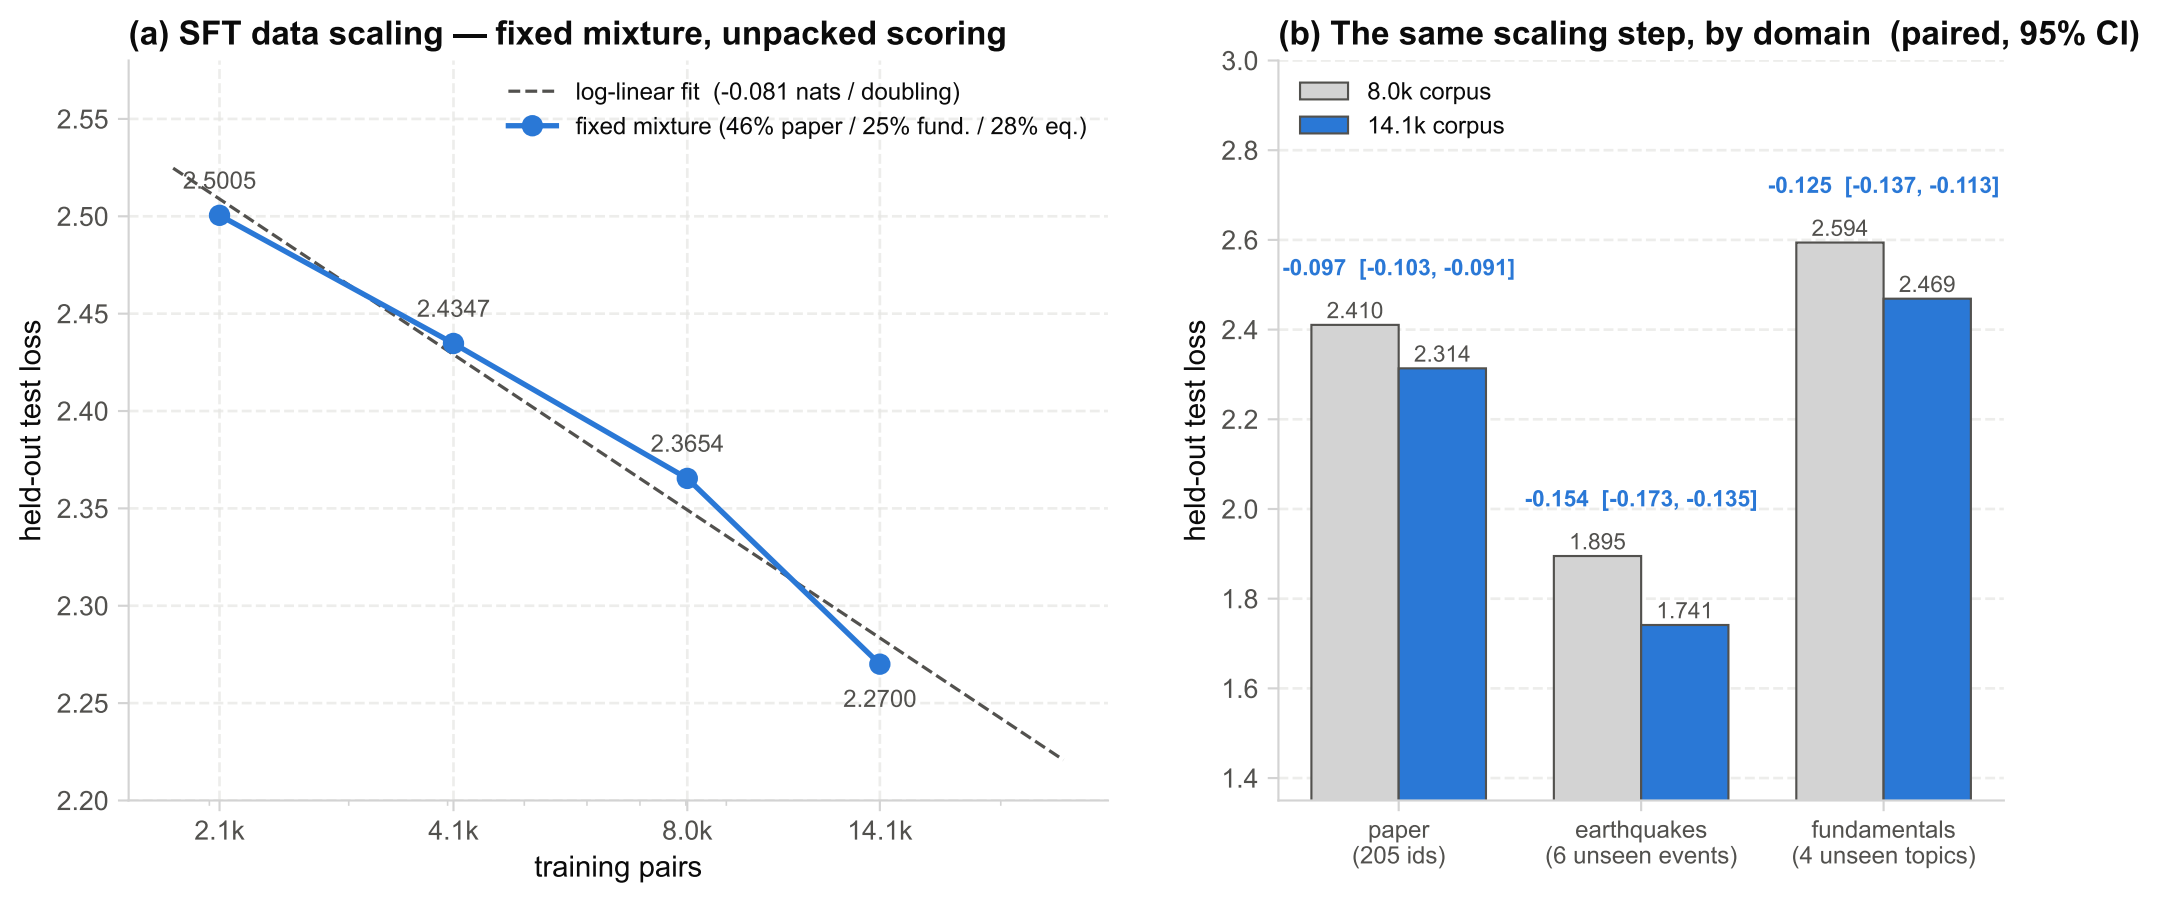
(a) SFT data scaling — fixed mixture, unpacked scoring
log-linear fit (-0.081 nats / doubling)
fixed mixture (46% paper / 25% fund. / 28% eq.)
held-out test loss
training pairs
2.5005
2.4347
2.3654
2.2700
2.1k
4.1k
8.0k
14.1k
(b) The same scaling step, by domain (paired, 95% CI)
8.0k corpus
14.1k corpus
-0.097 [-0.103, -0.091]
-0.154 [-0.173, -0.135]
-0.125 [-0.137, -0.113]
2.410
2.314
1.895
1.741
2.594
2.469
paper
(205 ids)
earthquakes
(6 unseen events)
fundamentals
(4 unseen topics)