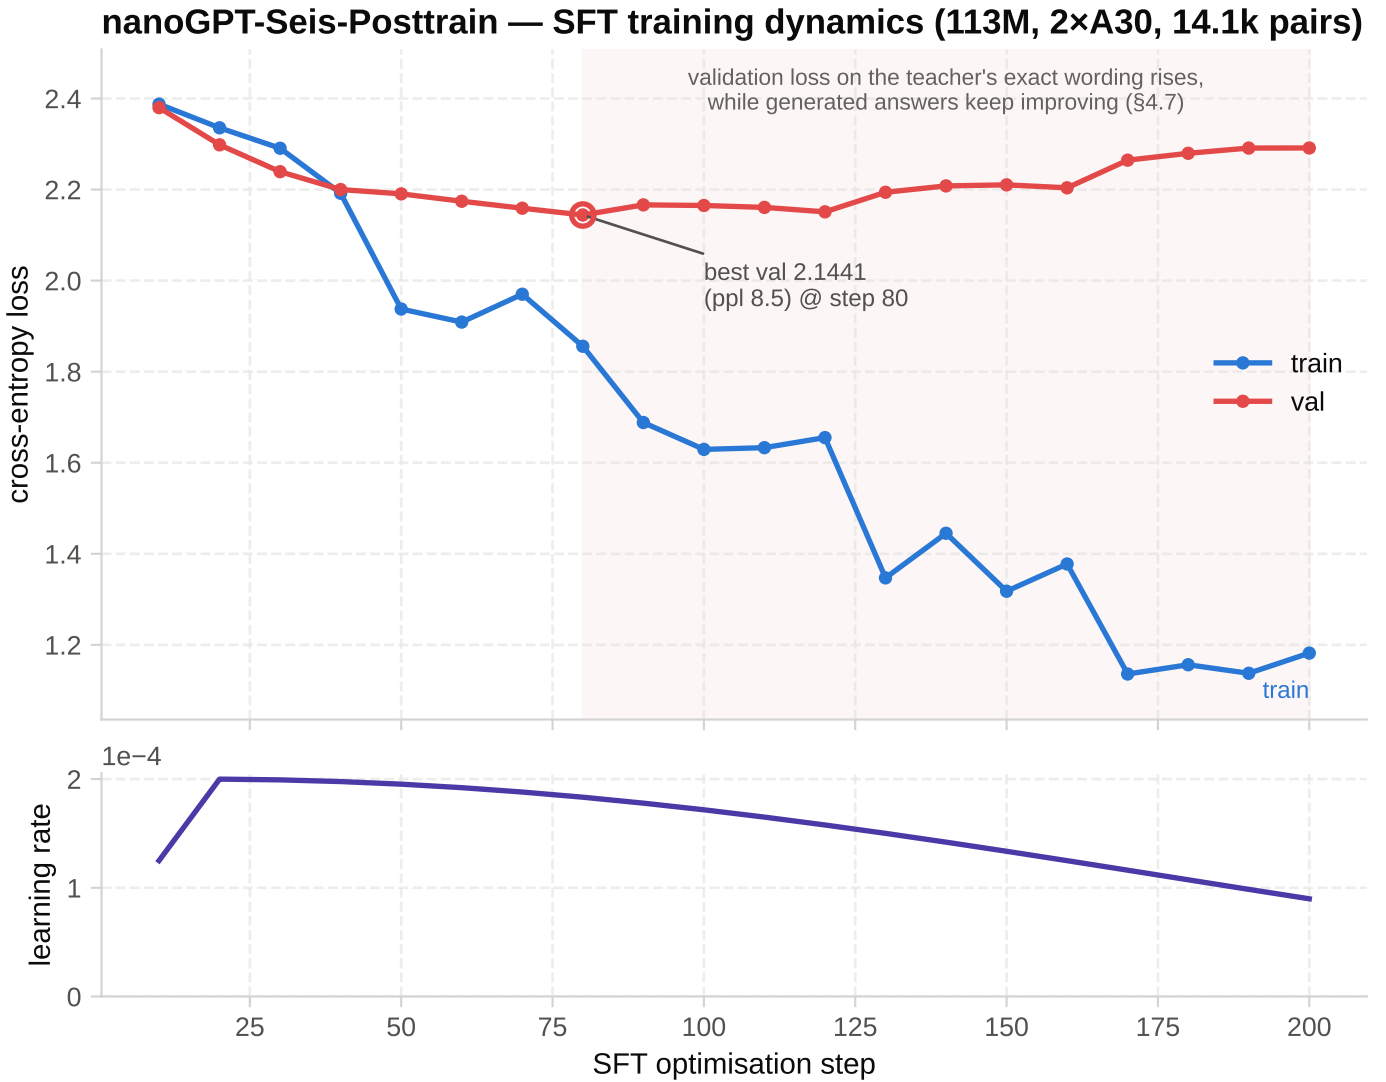
nanoGPT-Seis-Posttrain — SFT training dynamics (113M, 2×A30, 14.1k pairs)
validation loss on the teacher's exact wording rises,
while generated answers keep improving (§4.7)
best val 2.1441
(ppl 8.5) @ step 80
train
val
train
cross-entropy loss
2.4
2.2
2.0
1.8
1.6
1.4
1.2
1e−4
learning rate
2
1
0
25
50
75
100
125
150
175
200
SFT optimisation step