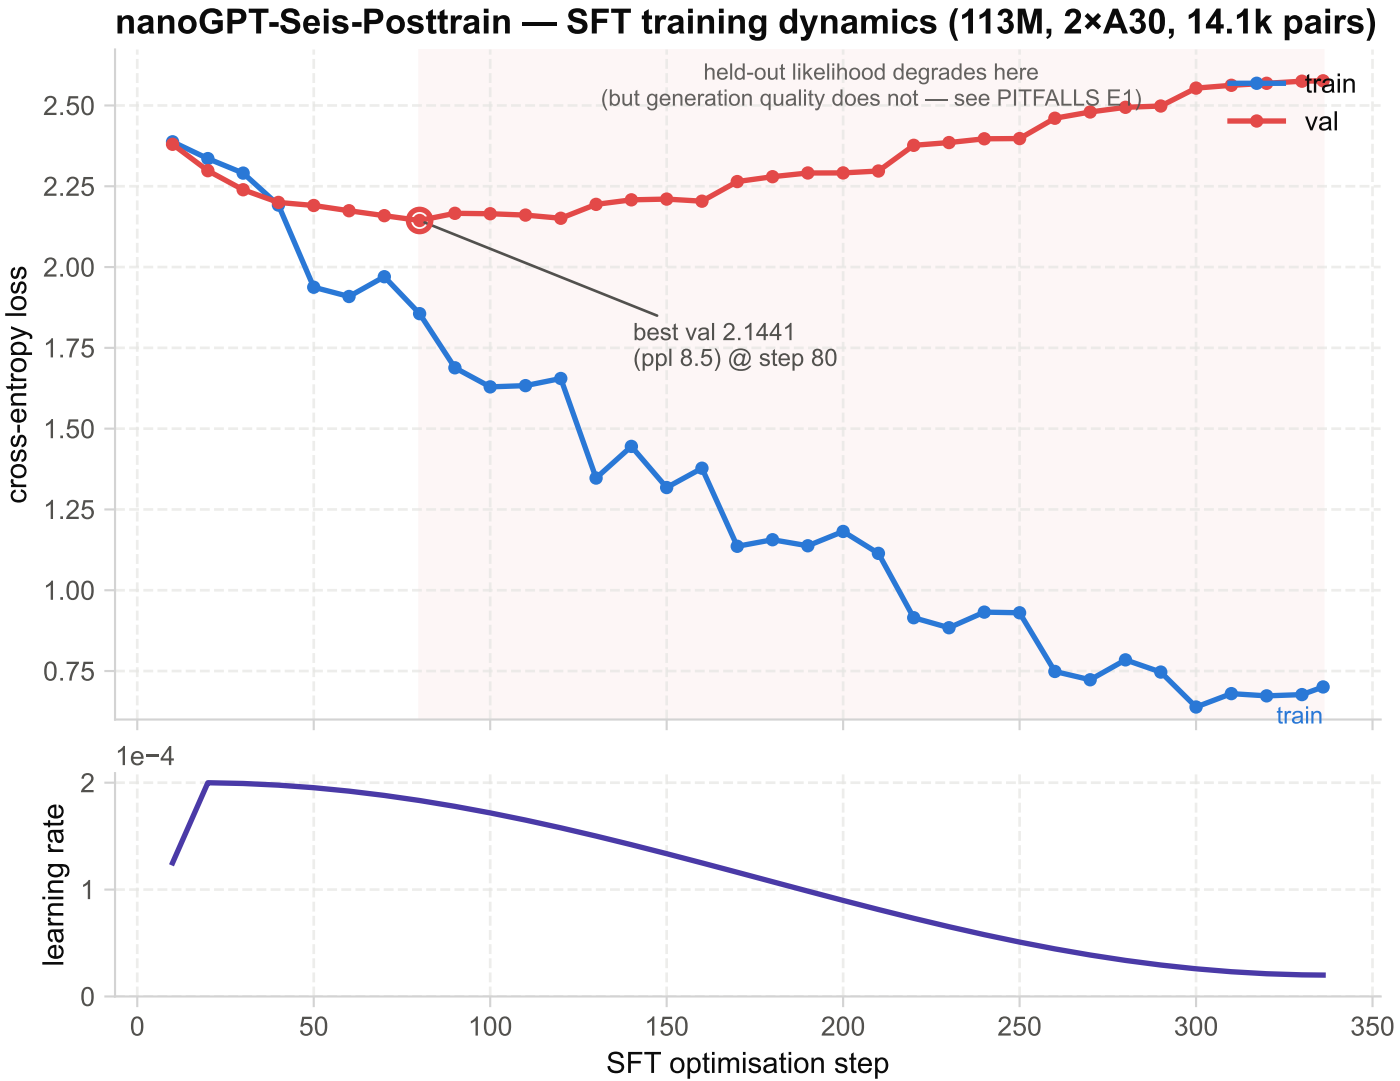
nanoGPT-Seis-Posttrain — SFT training dynamics (113M, 2×A30, 14.1k pairs)
held-out likelihood degrades here
(but generation quality does not — see PITFALLS E1)
train
val
best val 2.1441
(ppl 8.5) @ step 80
train
cross-entropy loss
2.50
2.25
2.00
1.75
1.50
1.25
1.00
0.75
1e−4
learning rate
2
1
0
0
50
100
150
200
250
300
350
SFT optimisation step